Lecture 1 - Linear Regression

Model Selection

When 2 models both fit the training data equally well, which one should we choose?

- use validation data and choose the better one
  - 一般来说 train : valid : test = 8:1:1
  - 一般over-fit时error on validation data会上升
- choose a simpler one (if without validation data)
  - Occam's Razor(奥卡姆剃刀准则)：如无必要，勿增实体
  - Inductive bias

Bias-Variance Decomposition

- D is a training set and $D = \{(x_1, y_1), \ldots, (x_n, y_n)\}$ $(D \sim P(D))$
  - 也就是说用由分布$P(D)$采样出来的D训练一个模型$f(x; D)$
- 对于一个测试数据(x, y)，平方误差的期望为 $E_D[f(x; D) - y]^2$
  - which means average error of $f(x; D)$ over infinitely many $D \sim P(D)$
  - we define $\bar{f}(x) := E_D[f(x; D) - y]$
  - now we have:

$$\begin{aligned} E_D[f(x;D) - y]^2 &= E_D[f(x;D) - \bar{f}(x) + \bar{f}(x) - y]^2 \\ &= E_D[f(x;D) - \bar{f}(x)]^2 + 2E_D[f(x;D) - \bar{f}(x)][\bar{f}(x) - y] + E_D[\bar{f}(x) - y]^2 \\ &= E_D[f(x;D) - \bar{f}(x)]^2 + (\bar{f}(x) - y)^2 \\ &= \text{Variance} + \text{Bias}^2 \\ & (E_D[f(x;D) - \bar{f}(x)] = 0) \end{aligned}$$

  - Variance: error caused by over-fitting particular D can be decreased by increasing the size *n* of D
    简单来说即是测试的预测与平均预测之差的平方
  - Bias: error caused by the model's ability to fit the data, or the model is not complex enough
    即平均预测与标签也就是实际值之差
- Trade-off
  - simple model: high bias, low variance
  - complex model: low bias, high variance

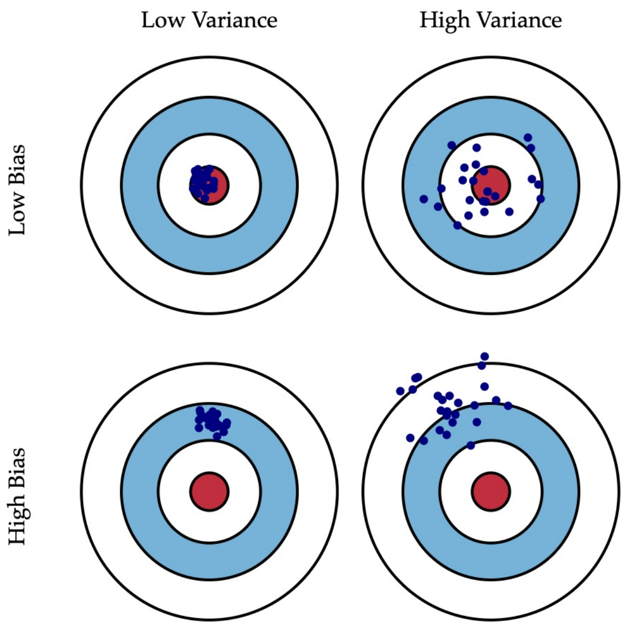

- One more thing
    - $\bar{f}(x) = E_D(f(x;D))$ is an average model over infinitely sample small D
    - $\hat{f}(x) = f(x;D,|D| \to \infty)$ is a model trained over an infinitely large D
    - 虽然上述推导使用了第一个函数，但在实际过程中数据自然越多越好，那么也就是第二个模型往往表现更好
    - example:

$$P(D) = \begin{cases} 1/3 & \{(-1,1),(0,0)\} \\ 1/3 & \{(1,1),(0,0)\} \\ 1/3 & \{(-1,1),(1,1)\} \end{cases}$$

        - $\bar{f}(x) = 1/3, variance > 0, bias^2 = 1$
        - $\hat{f}(x) = 2/3, variance = 0, bias^2 = 2/3$

Linear Regression

Basic Model

- D is a training set and $D = \{(x_1,y_1),\dots,(x_n,y_n)\}$. $x_i \in \mathbb{R}^d,\ y_i \in \mathbb{R},\ i = 1,2,\dots,n$
  We are going to train a model $y = w^Tx + b,\ w \in \mathbb{R}^d,\ b \in \mathbb{R}$
- How to determine w,b
    - *ERM(Empirical Risk Minimization)*
- Loss function
    - **squared loss** $(f(x_i) - y_i)^2$
- Objective:

$$\begin{aligned} &\min_{w,b} \sum_{i=1}^{n}(f(x_i) - y_i)^2 \\ =&\min_{w,b} \sum_{i=1}^{n}(w^Tx_i + b - y_i)^2 := L(w,b) \end{aligned}$$

- Gradient:

$$\frac{\partial L}{\partial w} = 2\sum_{i=1}^{n}(w^Tx_i + b - y_i)x_i$$
$$\frac{\partial L}{\partial b} = 2\sum_{i=1}^{n}(w^Tx_i + b - y_i)$$

  注意这里是一个向量对标量求梯度，结果还是一个向量
- Gradient Descent

$$w \leftarrow w - \alpha\frac{\partial L}{\partial w}$$
$$b \leftarrow b - \alpha\frac{\partial L}{\partial b}$$

  α is the **learning rate** and it's a **hyperparameter**
- 常用公式(已知$w \in \mathbb{R}^d$)

$$\frac{\partial w^Tw}{\partial w} = 2w$$
$$\frac{\partial w^Tx}{\partial x} = w$$
$$\frac{\partial x^TAx}{\partial x} = (A + A^T)x$$

Least Squares

- Least squares actually has closed-form solution! (which also regards as **Normal Equation**)
- we need some vectorized variables:

$$x = \begin{bmatrix} x_1^T, 1 \\ x_2^T, 1 \\ \dots \\ x_n^T, 1 \end{bmatrix} \in \mathbb{R}^{n\times(d+1)}$$
$$\hat{w} = \begin{bmatrix} w \\ b \end{bmatrix} \in \mathbb{R}^{d+1}$$
$$y = \begin{bmatrix} y_1 \\ \dots \\ y_n \end{bmatrix} \in \mathbb{R}^{n}$$

- the loss function now looks like this:

$$L(\hat{w}) = (y - x\hat{w})^T(y - x\hat{w})$$

  because we know that:

$$x\hat{w} = \begin{bmatrix} f(x_1) \\ \dots \\ f(x_n) \end{bmatrix}$$

- to compute $\hat{w}$ we want the gradient to be 0

$$\begin{aligned} \frac{\partial L(\hat{w})}{\partial \hat{w}} &= \frac{\partial[(y-x\hat{w})^T(y-x\hat{w})]}{\partial \hat{w}} \\ &= \frac{\partial(y^Ty - y^Tx\hat{w} - \hat{w}^Tx^Ty + \hat{w}^Tx^Tx\hat{w})}{\partial \hat{w}} \\ &= \frac{\partial(y^Ty - 2y^Tx\hat{w} + \hat{w}^Tx^Tx\hat{w})}{\partial \hat{w}} \\ &= -2x^Ty + 2x^Tx\hat{w} = 0 \end{aligned}$$

- now we focus on x^Tx:
  - 我们知道x^Tx是一个实对称矩阵，那么它是半正定的，也就是说该矩阵的特征值都是大于等于0的
  - 如果x^Tx不满秩，那么显然这个方程里$\hat{w}$是可能有多于一个解的，造成这种结果的可能原因是：
    - $d + 1 > n$:
      - $rank(x^Tx) = rank(x) \le \ min\{d+1, n\}$
    - x has repeated columns or rows
  - 如果x^Tx满秩，则特征值都大于0，我们可以对其进行特征值分解：

$$\begin{aligned} x^Tx &= U\Lambda U^T \\ &= U \begin{bmatrix} \lambda_1 & & \\ & \dots & \\ & & \lambda_{d+1} \end{bmatrix} U^T \\ (x^Tx)^{-1} &= U\Lambda^{-1}U^T \\ &= U \begin{bmatrix} \lambda_1^{-1} & & \\ & \dots & \\ & & \lambda_{d+1}^{-1} \end{bmatrix} U^T \end{aligned}$$

    - 那么这个方程有**唯一解** $\hat{w} = (x^Tx)^{-1}x^Ty$
    - 但依然有一个问题，就是λ_{d+1}可能很接近0，也就是x^Tx很接近一个非奇异矩阵，这就导致λ_{d+1}^{-1}会相当大，因此会导致计算上的问题(**numerical issues and instability**)

L2 Regularization

- Linear Regression with L2 Regularization (Ridge Regression):

$$minL(\hat{w}) + \lambda||\hat{w}||_2^2$$

  - λ is a positive **constant hyperparameter** which can control the strength of regularization (weight decay)
  - 简单来说就是避免某个w特别突出，使所有的w更平均一些
- loss function with L2 Regularization

$$J(\hat{w}) = (y - x\hat{w})^T(y - x\hat{w}) + \lambda w^T w \ \ (\lambda > 0)$$

- compute the gradient:

$$\begin{aligned}\frac{\partial J(\hat{w})}{\partial \hat{w}} &= -2x^T y + 2x^T x\hat{w} + 2\lambda\hat{w} \\ &= -2x^T y + 2(x^T x + \lambda I)\hat{w} = 0\end{aligned}$$

- 这是我们再对$x^T x + \lambda I$进行特征值分解，就可以得到:

$$\begin{aligned}x^T x + \lambda I &= U(\Lambda + \lambda I)U^T \\ &= U \begin{bmatrix}\lambda_1 + \lambda & & \\ & \dots & \\ & & \lambda_{d+1} + \lambda\end{bmatrix} U^T\end{aligned}$$

  so $x^T x + \lambda I$ is always **non-singular** and $\hat{w}$ is a unique solution!

$$\hat{w} = (x^T x + \lambda I)^{-1} x^T y$$

L1 Regularization

- Linear Regression with L1 Regularization (Lasso Regression):

$$\begin{aligned}&min\, L(\hat{w}) + \lambda||\hat{w}||_1 \\ = &min\, L(\hat{w}) + \lambda\sum_{i=1}^{d+1}|\hat{w}_i|\end{aligned}$$

  - induce sparsity of $\hat{w}$ (many dimensions of $\hat{w}$ will be 0)
  - Lasso stands for *Least Absolute Shrinkage and Selection Operator*
    - feature selection: 当参数很多时，用L1正则化可以筛选出真正有用的参数

Geometric view of Linear Regression

- Ideally, we want $X\hat{W} = y$
- 但是显然我们没办法找到一个model可以fit任何数据
- Consider the column space of X as $col(X)$, spanned by columns of X
  - $X\hat{W}$ lies in $col(X)$, but y may not
  - let $X\hat{W} = \hat{y}$, this $\hat{W}$ is the least square solution to $X\hat{W} = y$
  - Because $\hat{y}$ minimizes $||y - X\hat{W}||^2$, which is the object of Linear Regression, $\hat{y}(X\hat{W})$ should satisfy $y - \hat{y} \perp col(X) \leftrightarrow X^T(y - \hat{y}) = 0 \leftrightarrow X^T X\hat{W} = X^T y$