

# Ensemble Learning

- Combine (weak) base learners to obtain a stronger one
  - base learners should be as diverse as possible
- Bias-Variance Decomposition
  - High variance low bias: tree model
  - Low variance high bias: linear model

## Bagging

- **reduce variance**
- Suppose we can repeatedly sample  $D$  from  $P(D)$ , then let  $f(x) = \frac{1}{T} \sum_{t=1}^T f(x, D_t)$ 
  - $f(x) \rightarrow \bar{f}(x)$  when  $T \rightarrow \infty$ , then variance  $\rightarrow 0$
- However, we only have one  $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$

## Bootstrap

- Bootstrapping (自举) is a way that uses **random sampling with replacement**.
  - i. Sample  $n$  points from  $D$  with replacement
  - ii. Repeat i. for  $T$  times, get  $\hat{D}_1, \dots, \hat{D}_T$
  - iii.  $\hat{f}(x) = \frac{1}{T} \sum_{t=1}^T f(x, \hat{D}_t)$  as the bagged model
- 一共  $n$  个样本，有放回地每次选一个，一共选  $n$  次，样本在这  $n$  次中一次也没有被选到的概率为  $(1 - \frac{1}{n})^n \rightarrow \frac{1}{e} \approx 36.8\% (n \rightarrow \infty)$ ，因此有大约 36.8% 的数据不会被采样到
  - We can use the remaining data as **hold-out validation set (out-of-bag set)**
- Bootstrap has no theoretical guarantee. But empirically it can reduce test error greatly

## Random Forest

- the most successful bagging model
1. For  $t$  in  $[1, \dots, T]$ 
    - i. Sample  $n$  points  $\hat{D}_t$  from  $D$  with replacement
    - ii. Sample  $d' < d$  features  $F' \in F$  without replacement
    - iii. Build a full decision tree on  $\hat{D}_t, F'$  (can do some pruning to minimize out-of-bag error)
  2. End for
  3. Average all trees
- Simple, easy to implement, very powerful

## Boosting

- reduce bias (for linear models or tree models with certain height)

## Additive Model

- a general paradigm for *Boosting*

$$g(x) = \sum_{t=1}^T \alpha_t f_t(x)$$

1. Define  $g_t(x) = \sum_{i=1}^t \alpha_i f_i(x)$
2. At step  $t$ , keep  $g_{t-1}(x)$  fixed, then learn  $\alpha_t, f_t(x)$  by minimizing  $\sum_{i=1}^n L(y_i, g_{t-1}(x_i) + \alpha_t f_t(x_i))$
3. Finally,  $g(x) = g_T(x)$

## Adaboost

### General Form

- $D = \{(x_1, y_1), \dots, (x_n, y_n)\}, y \in \{-1, 1\}, x \in \mathbb{R}^d$

- Input:  $D$ , a weak learning algorithm  $A$

1. Initialize sample weights  $w_i^{(1)} = \bar{w}_i^{(1)} = \frac{1}{n}, \forall i \in \{1, \dots, n\}$
2. For  $t$  in  $[1, \dots, T]$ 
  - i. use  $D$  and  $\{w_i^{(1)}\}$  to train  $A$ , get a model  $f_t(x) : \mathbb{R}^d \rightarrow \{-1, 1\}$
  - ii. evaluate  $f_t(x)$ 's weighted classification error  $e_t$  on  $D$

$$e_t = \sum_{i=1}^n w_i^{(t)} \cdot 1(f_t(x_i) \neq y_i)$$

- iii. compute a weight  $\alpha_t$  for  $f_t(x)$

$$\alpha_t = \frac{1}{2} \log \frac{1 - e_t}{e_t}$$

$e_t \searrow, \alpha_t \nearrow$ . If  $e_t < 0.5, \alpha_t > 0$ ; if  $e_t > 0.5, \alpha_t < 0$  (如果这个模型错误率大于0.5, 显然直接结果取反就好了)

- iv. update sample weights  $\bar{w}_i^{t+1} = \bar{w}_i^t \cdot \exp(-\alpha_t y_i f_t(x_i))$ 
  - those  $y_i \neq \text{sign}(x_i f_t(x_i))$  get larger weights
- v. normalization

$$w_i^{t+1} = \frac{\bar{w}_i^{t+1}}{\sum_{j=1}^n \bar{w}_j^{t+1}}$$


by normalization, we have:

$$\sum_{i=1}^n w_i^{t+1} = 1$$

3. Combine  $T$  classifiers linearly

$$g(x) = \sum_{t=1}^T \alpha_t f_t(x)$$

- $g(x) \geq 0$ , predict  $+1$ , otherwise predict  $-1$



- 可以看出, 被分错的样本权重增加了, 而且多个分类器叠加后就能实现非线性划分的效果

## Adaboost Derivation

- Adaboost is an Additive model using exponential loss
  - $L(y, f(x)) = \exp(-yf(x))$
  - In Adaboost,  $L(y, f(x)) = \{e, \frac{1}{e}\}$
- At step  $t$ , we already have

$$g_{t-1}(x) = \sum_{j=1}^{t-1} \alpha_j f_j(x)$$

We need to optimize

$$\min_{\alpha_t, f_t} \sum_{i=1}^n \exp\left(-y_i (g_{t-1}(x_i) + \alpha_t f_t(x_i))\right)$$

Define

$$\begin{aligned} & \exp(-y_i g_{t-1}(x_i)) \\ &= \prod_{j=1}^{t-1} \exp(-y_i \alpha_j f_j(x_i)) \\ &= \bar{w}_i^{(t-1)} \cdot \exp(-y_i \alpha_{t-1} f_{t-1}(x_i)) \end{aligned}$$

to be  $\bar{w}_i^{(t)}$ , then we need to optimize

$$\min_{\alpha_t} \min_{f_t} \sum_{i=1}^n w_i^{(t)} \cdot \exp(-y_i \alpha_t f_t(x_i))$$

- In Adaboost, we just train  $f_t$  using its **original loss** and sample weights  $w_i^{(t)}$ , which means that we fix  $\alpha_t$  and optimize  $f_t$
- $\bar{w}_i^{(t)}$  和  $w_i^{(t)}$  只差一个常数
- After solving  $f_t$ , we need to optimize

$$\begin{aligned} & \sum_{i=1}^n w_i^{(t)} \exp(-y_i \alpha_t f_t(x_i)) \\ \Leftrightarrow & \min_{\alpha_t} \sum_{y_i=f_t(x_i)} w_i^{(t)} \exp(-y_i \alpha_t f_t(x_i)) + \sum_{y_i \neq f_t(x_i)} w_i^{(t)} \exp(-y_i \alpha_t f_t(x_i)) \\ \Leftrightarrow & \min_{\alpha_t} \sum_{y_i=f_t(x_i)} w_i^{(t)} \exp(-\alpha_t) + \sum_{y_i \neq f_t(x_i)} w_i^{(t)} \exp(\alpha_t) \\ \Leftrightarrow & \min_{\alpha_t} \exp(-\alpha_t) \sum_{i=1}^n w_i^{(t)} + (\exp(\alpha_t) - \exp(-\alpha_t)) \sum_{y_i \neq f_t(x_i)} w_i^{(t)} \\ \Leftrightarrow & \min_{\alpha_t} \exp(-\alpha_t) + (\exp(\alpha_t) - \exp(-\alpha_t)) e_t \end{aligned}$$

- Take gradient with  $\alpha_t$

$$-\exp(-\alpha_t) + (\exp(\alpha_t) - \exp(-\alpha_t)) e_t = 0$$

and we can get

$$\alpha_t = \frac{1}{2} \log \frac{1 - e_t}{e_t}$$

## Gradient Boosting Model

- For regression, no need for  $\alpha_t$ , because  $\alpha_t$  can be absorbed into  $f_t$  (因为都是实数)
- For squared loss, we need to optimize

$$\min_{f_t} (y_i - g_{t-1}(x_i) - f_t(x_i))^2$$

- we call  $r_i^{(t)} = y_i - g_{t-1}(x_i)$  **residual error**
- we iteratively train a new model  $f_t$  to fit the residuals  $\{(x_1, r_1^{(t)}), \dots, (x_n, r_n^{(t)})\}$
- Boosting Tree: Squared loss & Additive model &  $f_t$  is a tree
- Gradient Boosting Model
  - instead of fitting residuals, fit the negative gradient of loss, that is, let

$$r_i^{(t)} = - \left. \frac{\partial L(y_i, \hat{y})}{\partial \hat{y}} \right|_{\hat{y}=g_{t-1}(x_i)}$$

- essentially, use new models  $\alpha_t f_t$  to perform *gradient descent*
- for squared loss:  $L(y_i, \hat{y}) = \frac{1}{2} (y_i - \hat{y})^2$

$$- \frac{\partial L(y_i, \hat{y})}{\partial \hat{y}} = y_i - \hat{y}$$

## XGBoost

- perform second-order optimization & regularization (of tree complexity) & parallelization techniques